Reality First

Digital traces, automation, and an orthopraxy of judgment: a manifesto

Dimitri Diomaiuta

https://github.com/dda410/praxis

Contents

1 Preamble **3**

2 The problem **3**
- 2.1 The articulation tax . . . 3
- 2.2 Map drift . . . 3
- 2.3 The new temptation: wholesale delegation . . . 3

3 The two principles **4**
- 3.1 Principle I — the architecture: representations are answerable to traces . . . 4
- 3.2 Principle II — the allocation policy: one node is non-delegable . . . 4
- 3.3 Neither alone: why the practice is the conjunction . . . 5

4 The vocabulary **6**

5 The loop **6**
- 5.1 The shape . . . 6
- 5.2 The systems reading: a balancing loop governing a reinforcing one . . . 7
- 5.3 A phantom, caught: the loop working once . . . 7
- 5.4 What the law of requisite variety does and does not prove . . . 8

6 The division of labor **8**
- 6.1 The two machine seats . . . 8
- 6.2 The registered prior . . . 9
- 6.3 Delegate by authorship, not difficulty . . . 9
- 6.4 The oracle, refused a seat . . . 9

7 The artifact spectrum **10**
- 7.1 Axis one: bounded ↔ open-ended . . . 10
- 7.2 Axis two: well-modeled ↔ novel . . . 10
- 7.3 The live emergent document . . . 10
- 7.4 Roadmaps are plural and local; the spine is singular and map-less . . . 11
- 7.5 Evidence: one durable home for every trace . . . 11
- 7.6 Provenance: the audit graph . . . 12

8 The daily practice **12**

9 The five fragilities **13**

10 Scope limits **14**

11 Falsifiability **15**

12 Lineage **16**
- 12.1 The engine: representations answerable to reality . . . 16
- 12.2 The judgment node: why one act stays human . . . 17
- 12.3 The augmentation tradition: extending the witness . . . 18
- 12.4 The loyal opposition: kept because they push back . . . 18

13 Genus and name **19**
- 13.1 A practice, not a methodology . . . 19
- 13.2 The name . . . 20

14 The extraction **20**

15 Provenance of this manifesto **20**

1 Preamble

This is the manifesto of a working practice. It describes how to organize any work that *accumulates against a reality and must be answered for* — a research program, a long engineering effort, a commissioned delivery, a professional life — around two things: the **digital traces** the work leaves behind, and the **automation** (today, language models) that can absorb almost every cost of keeping those traces honest.

The practice in one paragraph:

> Organize the work as a self-correcting loop. Action makes contact with reality and leaves traces. Traces are the most authoritative account of what happened; every representation built on them — map, roadmap, plan, summary — is re-derived from the traces and disposable, and where representation and traces disagree, the traces win. A machine does all the friction around the loop: recording, distilling, surfacing, narrating, scaffolding — because that friction is what historically kills the practice. Exactly one node is never delegated: the situated human judgment that turns surfaced reality into a committed choice and signs for it. The representation's prescriptiveness is a dial, set by how bounded the work is.

Everything that follows introduces the concepts in that paragraph, grounds them, names where they are fragile, and states what would prove them wrong. The practice was not designed top-down; it was extracted from several working systems — a personal memory log kept across years, a data-science methodology, a compiler project, a bounded delivery against a brief agreed in writing with its recipient — that converged on the same design decisions independently. The manifesto distils that convergence, with its corrections applied rather than hidden.

2 The problem

Three failures motivate the practice. Each is old; the third is new.

2.1 The articulation tax

Faithful records of work die from friction, not from lack of conviction. The lab notebook decays into "we tried some stuff, it didn't work." The engineering log goes stale the week a deadline lands. The journal stops the day it becomes a chore. The cost of *articulating* what happened — writing it down, keeping it organized, keeping it findable — competes with the work itself, and the work wins. The valuable practice is not done badly; it stops being done at all.

2.2 Map drift

Every standing representation of work — the project plan, the backlog, the architecture document, the mental model — begins as a faithful account and then quietly stops being one. Reality moves; the map is groomed instead of re-derived; eventually the map is defended *against* reality. A plan that has drifted is worse than no plan, because it still carries authority. Whole traditions of failure live here: the master plan that flattens what it cannot see, the metric that replaces the goal it was meant to measure, the backlog item nobody can trace to an actual requirement.

2.3 The new temptation: wholesale delegation

Language models changed the economics of both failures at once. They can absorb the articulation tax almost entirely — and they can also, for the first time, plausibly absorb the *judgment*. The same machine that will happily maintain your log will happily write your reflections, propose your hypotheses, and make your decisions, fluently enough that the difference is invisible

on the page. So the old question — "how do we keep records honest?" — acquires a new, sharper companion: *what, exactly, must not be delegated, and why?*

The reality-first practice is one answer to both questions together.

3 The two principles

The practice rests on two principles. They are independent — neither can be derived from the other — and the practice is their *conjunction*. Getting their relationship right matters, because each alone collapses into something that already exists and does not need a manifesto.

3.1 Principle I — the architecture: representations are answerable to traces

Picture someone charting an unknown coast from a small boat.

The explorer's golden rule: *the chart is always redrawn from the coastline — never the coastline "corrected" to match the chart.* If the map says "bay here" and you sail up to solid cliff, you do not argue with the cliff; you cross out the bay and draw the cliff. The sea is the truth; the chart is disposable notes about it. A captain who trusts his beautiful old chart over the rocks in front of him drowns.

Stated precisely, the principle is an **authority ordering among representations**. Pure "reality" is never directly in hand — what is in hand are *traces*: the commit, the test result, the signed contract, the measurement. Traces are themselves representations — selected, encoded residues, not the territory. But they are the *least-derived* account of reality available, and the principle is that authority flows strictly down the derivation chain:

> reality → traces → derived representations (maps, plans, summaries)
>
> Each layer is answerable to the one above it, and is re-derived from it, never the reverse. Where layers disagree, the less-derived layer wins.

This is how the sciences already handle theory-laden data — data outranks theory even though data is never raw — and it is honest in a way the slogan "reality wins" is not. The trace is not the territory either; it is simply the closest thing to it that can sit in a file.

Three commitments follow directly:

- **Append-only traces.** History is never rewritten. A correction is a new trace that supersedes an old one visibly, not an edit that makes the past agree with the present.
- **Disposable representations.** Any map must be deletable and rebuildable from the traces. If it cannot be regenerated, it has become an authority of its own.
- **Re-derivation as a habit.** The map is not checked against reality when convenient; re-contact is a discipline with a cadence, because drift compounds in the gaps.

3.2 Principle II — the allocation policy: one node is non-delegable

The second principle governs the division of labor between the human and the machine, and it draws the line in an unexpected place: not at difficulty, but at **authorship-dependence**.

A simple test finds the line. Imagine the act performed by someone — or something — else, with a *provably identical* result, and ask whether you would take that deal.

- **The dishwasher.** Would you let a machine wash your dishes if they came out exactly as clean? Obviously yes. The clean dishes were the only thing you wanted; who produced them is irrelevant. The value is in the output → **delegate**, gladly, even though it is the bulk of the chore.

- **The marathon.** Would you pay someone to run your marathon for you and pin the medal on *you* — same finish time, your name in the results? Obviously not. The medal was never the point; running it was. An identical outcome leaves nothing of value behind. The value is in the doing → **keep** it.

Both are tedious, sweaty, hours of effort — so the line between them is *not difficulty*. In the dishes, the effort is the **price** you pay for the result; in the marathon, the effort **is** the result. The same test settles the cases this practice cares about: would you hire a ghostwriter for your wedding vows, even if they moved the room to tears just as much? No — vows are worthless unless they are yours. That "no" is the same "no" as letting a model author your reflection, your scientific judgment, or your sign-off.

The distinction is ancient. Aristotle separated *poiesis* — production, whose end is the product — from *praxis* — action, whose end is in the doing itself. The allocation policy in one line: **delegate the poiesis, keep the praxis.** A line of code is authorship-independent — any competent agent yields an equivalent line — so it is fully delegable, even when it is the bulk of the visible, skilled labour. The design choice and the sign-off are authorship-dependent, so they stay. The machine may write every line; the human decides what to build and signs that it was built right.

What stays human, concretely, is one node: **the situated judgment that metabolizes surfaced reality into a committed choice, and signs for it.** Signing is the act that constitutes you as the one answerable — and because you sign, you must inspect what you sign with distrust, deep inspection, and a critical spirit, as you would review external work that goes out under your name.

3.3 Neither alone: why the practice is the conjunction

Each principle, taken alone, names something that already exists.

- **The loop without the judgment node** is just cybernetics. A thermostat runs Principle I perfectly — re-contact reality, correct the representation — with no witness, no signature, no judgment. Most control systems do. Nothing new there.
- **The judgment without the loop** is just record-keeping. A signed affidavit is an answerable, authored account with no loop at all. Nothing new there either.

What is distinctive is the conjunction: a self-correcting loop between reality and a disposable representation, *in which the one node that is not automated is the situated human judgment* — while everything else is offloaded to a machine. And the two principles need each other in both directions:

- **Principle II selects the domain where Principle I matters.** The loop applies everywhere; the *practice* applies exactly where the value of the output depends on someone being answerable for it over time. Where nothing must be answered for, run the thermostat — full automation is correct, and this manifesto does not apply (see Scope limits).
- **Principle I is what makes Principle II's node trustworthy.** The judgment literature is blunt about this: intuitive expert judgment is reliable only in environments that provide valid cues and fast, honest feedback (Kahneman & Klein, 2009). An unlooped judge drifts exactly like an unlooped map. The loop — re-contact, re-derive, compare — is what *calibrates* the human comparator. The judge needs the loop as much as the loop needs the judge.

One more asymmetry, worth stating because it dates this document honestly: **Principle I is timeless; Principle II is era-dependent.** The direction of authority between map and territory does not move as machines improve. The *location of the fence* does — work that was

kept human in one capability regime becomes delegable in the next (writing code already crossed; proposing hypotheses may cross). Principle II is therefore not a fixed list of protected tasks but a **policy**: at any moment, delegate everything authorship-independent, however hard, and keep the choice and the signature. The policy is stable; its extension moves with the machines.

4 The vocabulary

The practice needs few terms, but it needs them exactly. Eight are load-bearing; everything else can be said in ordinary words.

World-trace the residue of *reality*: the commit, the signed contract, the test result, the meeting's actual outcome. The authoritative layer. World-traces have veto power over every derived representation, and a *changed* world-trace is the one thing allowed to interrupt you and pull the framing upstream.

Signal an incoming stimulus that is not yet a trace of anything: an email arrives, an idea is floated, a thought occurs. A signal changes nothing until you form intent and act on it. Treating every signal as authoritative floods the judgment with noise.

Work-trace the residue of *your activity*: the conversation you had, the decision you recorded, the dead end you explored. Provenance and lineage — narrative, not setpoint.

Comparator the human node of the loop: the one who holds the derived view against the traces, judges, chooses, and signs. The word is borrowed from control theory, with one correction noted below.

Setpoint what the loop regulates toward. For the account itself: "the representation faithfully matches the traces." For bounded work: the deadline and the contract — which are themselves world-traces.

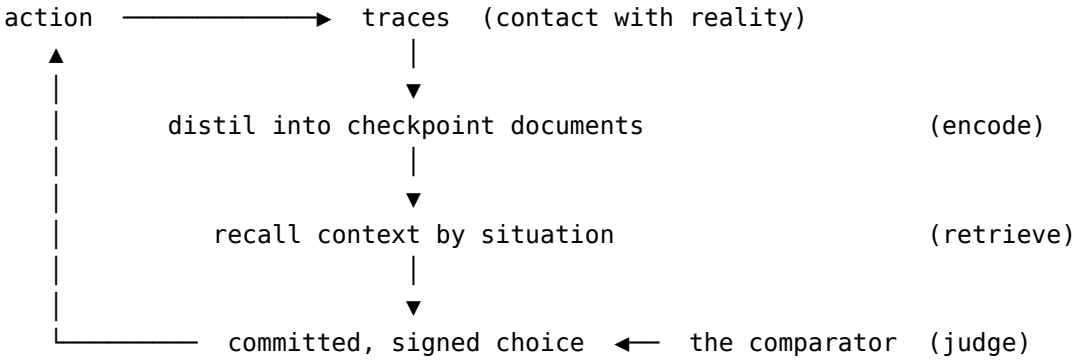
Phantom an item in a derived representation whose provenance reaches no world-trace. A requirement nobody actually imposed, a belief with no residue behind it. Phantoms are the characteristic failure of standing maps, and the provenance discipline below exists to make them visible: they are the nodes with no root.

The spine the one append-only narrative log of the practitioner across all work — identity, durable learnings, accumulated judgment. Singular, map-less, and never rewritten.

Derived view (or *roadmap*) any disposable representation re-derived from traces: a project roadmap, a status summary, an index. Plural, local, and deletable by construction.

5 The loop

5.1 The shape



Action leaves traces. Traces are *encoded*: distilled into checkpoint documents — not chronological records but crystallisations, engineered for a dual audience, serving at once as the human's resumption point and as portable context that primes a model to reason over the whole picture.

Checkpoints are *retrieved* by situation, not loaded wholesale: the project you are in, the threads that are fresh, the horizon you are on cue what comes back into view. (This design is modeled on the encoding-specificity findings of the memory literature — retrieval works best when the cues at recall match the cues at encoding — as a guiding analogy, not a claimed implementation.) Retrieval informs the *judgment* — the non-delegable node — which issues new action, which leaves new traces.

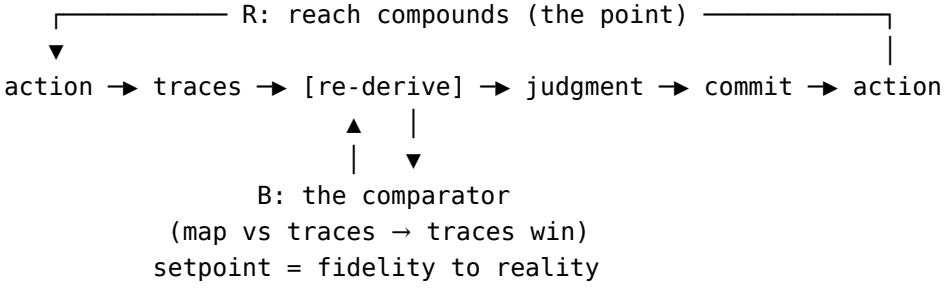
The machine operates on every arc except the judgment: it distils, it recalls, it narrates. It does not judge. And the loop re-contacts reality on every pass — which is why the account cannot quietly drift into elaborating its own prior plan.

5.2 The systems reading: a balancing loop governing a reinforcing one

In systems-dynamics vocabulary (Meadows), the practice is two coupled loops, and the whole design is in the coupling.

The *governing* loop is **balancing** — textbook negative feedback. Stock: the fidelity of the representation to the traces. Setpoint: "the account matches reality." Error signal: the gap between map and traces. Corrective flow: traces win, the map is pulled back.

The loop *worth running* is **reinforcing**. Define **reach** as the span of reality a single practitioner can faithfully account for — across time and across parallel threads. Each honest cycle enlarges the account, which improves retrieval, which sharpens judgment, which produces more faithful action, which enlarges the account again. Reach compounds. That compounding is the point of the practice — it is what "extending human reach" means, mechanically.



The load-bearing insight is pure systems theory: **a reinforcing loop left ungoverned runs away.** Without the balancing loop, reach does not compound fidelity — it compounds *drift*: the account keeps growing, but grows away from reality, larger and less true. The practice is a designed inverse of the classic "limits to growth" trap: instead of letting the reinforcing engine slam into an external limit, you install the balancing loop internally, on purpose, as a governor. **R is why you bother; B is what keeps R honest.**

Two of the practice's habits are, in these terms, interventions at the highest-leverage points:

- **Re-derive from traces every session** is *shrinking the perception delay.* Delays in balancing loops cause overshoot and oscillation; a map that is not resampled for weeks accumulates error silently.
- **Traces win over the map** is the anti-Goodhart rule: regulate against the least-derived account available, never against the representation you are fond of. Balancing loops fail precisely when they control against the believed state instead of the actual one.

5.3 A phantom, caught: the loop working once

One real incident shows the whole mechanism. A bounded delivery against an agreed brief carried a deliverable — a separate dossier of source documents — that had entered as a *self-imposed idea* in a working session and been encoded into the roadmap *as if the brief required*

it. It survived there for weeks, untethered. The fix was exactly the loop: a session-opening re-contact with the actual world-trace (the brief, agreed in writing with its recipient before the work began) revealed the map had grown a requirement reality never imposed. The item was struck through — not deleted — with the date and the reason, so the correction is itself a trace.

Every element of the practice appears in that story: the phantom (an item whose provenance reached no world-trace), the perception delay (weeks without resampling), the corrective habit (verify the real state at session opening), and the append-only repair (strikethrough, never erasure). Drift is not hypothetical; it is the standing failure mode of every prescriptive document, and the only defence is the habit of re-deriving from traces.

5.4 What the law of requisite variety does and does not prove

It is tempting to ground the human node in Ashby's law of requisite variety — a regulator can control a system only if it commands at least as much variety as the disturbances it must absorb — and conclude that the comparator *must* be human, as if by theorem. The law proves no such thing. It is completely neutral about what kind of system the regulator is; if a machine attains requisite variety for a domain, Ashby *licenses* automating that domain's comparator.

What the law actually supplies is the **criterion the fence answers to**: the comparator must command variety matching the reality it judges — including everything tacit, situated, and contextual that never made it into the traces. The claim that, for answerable work, only the situated human currently meets that criterion is an *empirical, dated* claim — supported today by what we know of tacit knowledge (Polanyi) and expert judgment (Dreyfus), and honestly subject to revision. This is exactly why Principle II is a policy rather than a theorem, why the fence is held as a norm rather than a mechanism, and why its location moves between capability regimes while the architecture of Principle I does not.

6 The division of labor

6.1 The two machine seats

The machine sits in two distinct seats, on either side of the human node.

- **The interlocutor** operates on the *forming* of the decision: divergent first (widen the view), then convergent (narrow it), then adversarial (attack the assumptions). Properly understood it is a **stimulator**: the diverging and converging happen in the human's head; the model is the provocation that widens the aperture and the foil that surfaces hidden assumptions — a stimulus to cognition, never its replacement.
- **The chronicler** operates on the *recording* of the work: making dense traces — code, choices, dead ends — legible; translating technical work into a followable story; re-rendering the derived views. The model is the chronicler, not the author.

The human's non-delegable acts sit precisely *between* the two seats: the model helps you form intent and helps you record what happened; it is fenced out of the instant where you actually choose and the instant where you actually witness.

The seats also give the practice a temporal shape. At *inception*, the interlocutor earns its keep: one or more sessions that diverge across the problem space and converge on a crystallised problem statement and a roadmap. In *execution*, the chronicler and the loop take over: recording, recalling, re-projecting, with the roadmap kept honest by re-derivation. The choice brackets the front; the sign-off brackets the end.

6.2 The registered prior

The interlocutor seat carries a danger the early versions of this practice underestimated: **the menu is the decision.** Decades of judgment research (anchoring; choice architecture) show that whoever generates the option set substantially steers the choice, even when the chooser formally decides. A model that frames your options has done much of your deciding, invisibly, while leaving the signature to you. The fence around the final pick does not protect the agenda.

The countermeasure is a discipline: **write your own position down before invoking the model.** State what you currently believe, what you expect, what you would do absent advice — then open the dialogue. The stimulator then widens an existing aperture instead of framing an empty one, and you can see afterwards exactly where it moved you. Ask for the case *against* your registered position before asking for options. A prior you never wrote down is a prior the model writes for you.

6.3 Delegate by authorship, not difficulty

Everything authorship-independent is handed to the machine eagerly — and "everything" is meant literally, including the bulk of the visible, skilled labour. Surfacing, orienting, indexing, formatting, summarizing-for-orientation, scaffolding, drafting, and in software the writing of nearly every line of code: all of it is poiesis, valued by its output, delegable in full. ~~Partial delegation of the delegable layer is not caution; it is a babysitting tax that frees no one.~~ *(Struck 2026-08-30: true of delegation aimed at the output; not of delegable work retained to calibrate the judge — see fragility five.)*

What is kept is praxis: the choice of what to build, the judgment against reality, the reflection, the signature. Kept not because the machine cannot produce convincing versions of them — it increasingly can — but because a convincing version produced by a non-witness is worth nothing. The account derives its value from who answers for it.

6.4 The oracle, refused a seat

The model invites a third use, distinct from both seats: the **oracle**. The two seats ask it to work on what you already command — you can judge the output, so delegation is safe by construction. The oracle use asks it what you do *not* know, and there the danger is routinely misdiagnosed. It is not that asking is wrong, nor even that the model may be wrong. It is that an oracle's answer is a derived representation whose provenance reaches no world-trace — and the moment you act on it unchecked, it becomes authoritative: a phantom at birth. Oracle use is the map-first inversion re-entering through the window — a non-answerable representation outranking reality — inside the very practice built to prevent it.

The operative question is therefore never "is the model reliable here?" — that question has no honest answer at the moment you need one. The question is: *does this answer have a path to reality before I sign on it?* Two zones follow:

- **Convertible.** You cannot judge the answer directly, but you can make it touch reality before signing: run the code, check that the cited source exists and says what is claimed, do the experiment. The loop closes through traces, and the answer enters the system the only way anything may — as a trace of contact, not as an utterance taken on authority.
- **Unclosable.** No affordable contact with reality stands between the answer and the action. The loop cannot close, and the model's representation quietly takes the setpoint's place.

The rule in one line: **a model's answer is a signal, never a trace.** It may provoke intent like any other signal; it earns authority only through contact with reality. Where no such contact

is affordable, the honest moves are to record the ground as what it is — "acting on the model's unverified claim, because X" — or not to act at all.

7 The artifact spectrum

The practice does not prescribe one artifact. It prescribes a **dial**, set by two axes of the work.

7.1 Axis one: bounded ↔ open-ended

How delivery-bounded is the work?

- **Bounded delivery** — a defined end-state, a deadline, phases, something to sign. Here the right artifact is the **live emergent document** (below): a roadmap that is descriptive and prescriptive at once, re-derived from the traces, driving toward an end.
- **Open-ended accumulation** — no end-state, no deliverable; identity and judgment forming across years. Here a standing prescriptive map is exactly the liability to refuse: the right artifact is the pure append-only **spine**, with views derived only at review time and then discarded.
- **Between** — work that is open-ended overall but bounded per chapter (a research program: questions arise indefinitely, but each hypothesis has a falsifiable end) maintains a map scoped to the bounded unit and nothing wider.

The artifact's *prescriptiveness* slides along this axis; the loop and the judgment node are invariant across the whole range. A methodology is not better or worse in the abstract — it shines on the problem it was shaped to and dims when lifted out of it. Choose the artifact by the shape of the work, not by allegiance to a method.

7.2 Axis two: well-modeled ↔ novel

There is an honest counter-tradition this practice must answer rather than ignore: in some regimes, *plan-primacy wins*. Aviation checklists, building codes, surgical protocols — standing prescriptive maps that systematically outperform situated judgment in their domains, with strong evidence behind them. Reality-first is not a universal law.

The resolution is a second axis: **how well-modeled is the domain?** Where the domain model is mature, failure modes are enumerated, and novelty is low, the map has earned authority — follow the checklist, and let deviation from it be the thing that requires judgment and a signature. Where the domain model is poor, the work is novel, or the territory is still being charted, the map cannot have earned that authority, and reality-first is the only honest stance. Most knowledge work — research, design, writing, early engineering — lives at the novel end. But the practice should say plainly: the explorer redraws the chart *because the coast is unknown*. Once a coast is charted to death, you follow the chart, and that too is fidelity to reality.

7.3 The live emergent document

For bounded work, the central artifact is a single living roadmap with three properties:

- **Descriptive and prescriptive at once.** It mirrors the current state of the project, re-rendered from the traces each time it is touched — and precisely *because* it is current, it can legitimately say what to do next. It is prescriptive only for as long as it stays honestly descriptive; the moment it decrees rather than reflects, drift has crept in.
- **Emergent.** It is not authored top-down as a plan; it condenses out of the accumulating traces and shifts as they shift.

- **Disposable.** It is a regenerable view: delete it and rebuild it from the traces. That deletability is the architectural guarantee that it never becomes the thing you are loyal to.

Under delivery pressure the clean separation of "log upstream, map downstream" collapses into a single file carrying three strata at once, and this is correct, not a compromise:

- an **append-only decision ledger** — each shipped item bound to its evidence (a commit hash, a date, a rationale), each scope decision frozen with its date;
- **corrections that preserve their own history** — a superseded plan is struck through with the date and the reason, never deleted, so the revision is itself testimony;
- a **live forward projection** — the phase breakdown, the verified current state, the "remains to be done."

7.4 Roadmaps are plural and local; the spine is singular and map-less

The granularity rule: **one balancing loop → one roadmap.** The unit is not "project" or "team" but *the unit with a single setpoint and a single signer*. Most projects are one deadline-shaped setpoint and get one roadmap. A sub-part earns its own roadmap only when it is itself a bounded sub-delivery with its own falsifiable end and its own owner; a mere phase of the same delivery is a section, not a new map. Where the horizon splits is where a new roadmap buds.

The placement rule: **roadmaps live next to the traces they are derived from** — in the project's own repository or workspace — never centralised in an orchestrator. Re-derivation must be local and cheap; "traces win" must be a low-latency check; and no central store of maps should ever compete with the spine for authority. The spine — the one open-ended log of the practitioner across all work — holds only what no single workspace can hold: identity, durable learnings distilled up from closed projects, and a *derived* index of what is fresh, active, and dormant. The spine **references, never duplicates**: it finds the roadmaps; it does not keep them. When a project closes, its roadmap is disposable — distil the durable learnings up into the spine and let the map dissolve. The reach stays; the scaffolding goes.

7.5 Evidence: one durable home for every trace

Every world-trace a decision rests on needs *one authoritative, re-contactable home*, so the derived views can be checked against it. Three tiers, by the durability of the referent:

1. **Durable and addressable → reference it.** A commit, a file, the signed contract, a stable ticket. Point at it; a copy would start drifting from the original.
2. **Ephemeral or external → capture it.** Email threads, screenshots, saved PDFs, pasted excerpts — into a per-project `traces/` (or `evidence/`) folder, each item carrying when it was captured and which decision it grounds. Here the copy *is* the faithful move: it does not drift from a living referent, it prevents the referent from vanishing.
3. **No footprint at all → externalise at the intent statement.** A hallway conversation, a verbal commitment, a realisation: write it down as you commit — "doing X because Y told me Z, <date>." That line becomes the world-trace.

This is deliberate, scoped capture by the comparator of the specific evidence a decision rests on — the opposite of ambient ingestion. Pulled deliberately and scoped to a decision: yes. Pushed automatically as an inbox: refused. And the capture discipline has a scope rule: capture what is *decision-relevant*, not everything — the cure for forgetting is better cue-matching, not more storage.

One case deserves naming rather than faking: some grounds are **tacit**. You know more than you can tell, and some judgments rest on experience that has no externalisable footprint. The honest move is to record the tacit ground *as such* — "doing X on my judgment, from experience

with similar cases" — rather than fabricating a documentary basis. That node is un-auditable by design, because it *is* the witnessing act. Recording it honestly is better provenance than pretending.

7.6 Provenance: the audit graph

Every artifact — a roadmap line, a todo, an idea, a decision — carries two links: a **down-link** to its origin (which world-trace or signal it came from) and an **up-link** to what it became (the intent or action it fed). The graph this closes is the drift-detector: an item whose down-link never reaches a world-trace is a phantom, flagged for pruning; a backlog item whose originating signal has gone stale is a candidate for review-time retirement.

And one asymmetry, which is Principle I in miniature: **world-traces may trigger re-entry; the backlog may not.** A failing test, a drifted metric, a re-read contract clause — a changed world-trace is allowed to interrupt and pull you upstream, because it is the setpoint. Todos and ideas merely *surface* when you orient, and wait to be chosen. A backlog that fires on its own is a map decreeing action — the exact inversion the practice exists to prevent. Intent is **formed at the confluence** of verified reality, surfaced signals, and consulted candidates; it is never *derived from* any artifact. The moment a todo "generates" the next action, the comparator's job has been quietly handed to the map.

8 The daily practice

A session is a cognitive block, not a day. Inside it you walk the loop once; the two human bookends are the non-delegable node.

1. **Orient — re-contact reality first.** Before trusting any document, resample the world: the log, the files, the test counts, what the contract actually says. Diff that against the roadmap; where they disagree, the traces win. This is the balancing correction firing at session start — minimizing the perception delay, the single highest-value habit in the practice.
2. **Surface — by situation, not everything.** Pull into view only what this moment cues: this project, fresh threads, the relevant horizon.
3. **Form intent — the comparator, with the stimulator.** Write your registered prior. Then let the model widen, narrow, and attack it. Then *you* commit to a choice, sign it, and record its provenance — which world-traces it regulates against, which signals and candidates it consulted, which grounds are tacit.
4. **Act — the only writer.** Do the work; hand the machine every authorship-independent piece. Action leaves traces. Nothing enters the system's state except through action.
5. **Chronicle — the machine distils.** It compresses the session's dense traces into legible narrative and re-renders the derived views. You sign the chronicle.
6. **Reflect — the comparator again.** Intent versus outcome, plus residue: what did this session leave behind that matters beyond it? Append to the spine; promote durable learnings out of the disposable map.

Rhythm: **orient → surface → intent → act → chronicle → reflect.** The friction in the middle belongs to the machine; the two ends — choosing and witnessing — are yours.

Two refinements:

- **The `contradict` move.** A standing use of the model whose only job is to surface *what contradicts the current framing* — never to maintain or polish it. This is the deliberate countermeasure to the model's own fluency (fragility three, below): an agent pointed at the cracks, because the default direction of a helpful model is to paper them over. The

verdict on a surfaced contradiction stays with the comparator — it is structurally safer to delegate *finding* a contradiction than to delegate *judging it resolved*.

- **Cadenced re-contact is legitimate for bounded work.** The purist "the human triggers everything" is right for the open-ended spine. But a deadline is itself a world-trace, and a scheduled re-orient is not ambient ingestion — it is shrinking the perception delay. The cadence resamples reality; the comparator still forms every intent. What stays forbidden is any automation that *decides* for you.

9 The five fragilities

The practice is load-bearing but brittle in five places. Naming them is part of the practice; each has a countermeasure, and none has a guarantee.

1. **The fence is a norm, not a mechanism.** Nothing technically stops you from rubber-stamping a machine-drafted reflection or accepting "the data suggests X" as your judgment. The moment you do, the artifact silently stops being yours while looking identical. There is no runtime check — only discipline, and the adversarial inspection that signing demands. As models grow more persuasive, the gradient across the fence steepens. The failure is invisible and self-inflicted.
2. **Some fences are thin.** Where the machine is plausibly capable of the protected act — proposing the hypothesis, ranking the options — the line between "orienting you" and "making the call" narrows. The discipline is to have the model *narrate options without ranking them* where the choice is the protected act, and to notice when orientation has quietly become recommendation.
3. **Fluency papers over the cracks.** The practice depends on reality being allowed to overturn the framing — and a model that is too good at making the existing plan feel coherent will smooth exactly the contradictions that should send you back upstream. The model's fluency is itself a risk to the practice's empiricism. Countermeasure: the `contradict` move, run as a habit, plus the session-opening re-derivation from traces — not cleverer maintenance.
4. **The menu is the decision.** Whoever generates the option set steers the choice, even when the human formally decides. The fence protects the final pick but not the agenda — and agenda-setting is where much of the steering power lives. Countermeasure: the registered prior, written before the model speaks, and the standing request for the case *against* it.
5. **The comparator drifts, and nothing in the loop catches it.** The loop as described recalibrates only the map: re-derive, diff, correct. No arc recalibrates the judge — yet the judge's variety comes from having done the work (Polanyi again: tacit knowledge is acquired in the doing, not in the watching), and total delegation of the doing silently switches off the feedback loop by which judgment is maintained. This is fragility one with a structural cause — eroded competence — in place of a moral one, and it is undetectable by construction: rubber-stamped output looks identical to inspected output, and the correction rate is ambiguous at both ends. A falling rate is ambiguous between "the machine improved" and "I stopped seeing." A high rate on a task recorded as delegated is ambiguous between "the machine is failing" and "I am doing the work through it and crediting it" — the Clever Hans effect, as Kambhampati (2024) observes of human-in-the-loop prompting: the horse that appeared to count while reading its handler, and the handler steering "even if they didn't set out to do so deliberately." The effect is defined by being unnoticed, so a tracked rate catches only its visible half; a back-prompt, a reworded ask, an "are you sure about step three?" is a correction round and is chronicled as one. The first ambiguity misattributes the machine's work to the judge; the second misattributes the judge's work to the machine, and it corrupts the evidence Principle II moves on — the

fence is moved outward on the experience that the machine now does the work. The test is not when the judge contributed but what the contribution carried: the line runs between a prompt that specifies the problem and a prompt that injects its solution, whenever either is issued. What has crossed is only what the machine added beyond what the prompts carried; a spec so complete that the machine only formats it delegates the formatting and nothing more. An act has crossed only when the machine's output passes a sound check the judge did not author the answer to, and the chronicle records what the prompts carried, not the word "delegated." (Kambhampati's closing point, that the arrangement becomes questionable when the human cannot verify the answer, is the unclosable zone of the oracle rule.) The better the delegation, the less informative the experience of supervising it (Bainbridge's irony of automation). The serious counterargument — the critic does not paint, the editor does not write novels — cuts the other way on inspection: those evaluators were formed by decades of feedback loops aimed at evaluation itself, with consequences for being wrong. Judgment is a stock maintained only by its own loop. Countermeasure: a *second balancing loop, pointed at the judge* — deliberate re-contact with the delegated layer (periodically do by hand what you normally delegate, not for the output but to resample your own competence: the inward analogue of the session-opening re-derivation); seeded errors, the blind proficiency test (deliberately defective output placed among the genuine, to measure whether detection still fires); and a tracked correction rate whose decline is treated as a signal to disambiguate, never as a success, and whose height is treated as a signal that the task has not crossed, never as delegation. The practice began by naming two things that decay — the map and the account. There is a third: the comparator. Each needs its own corrective loop.

10 Scope limits

The practice is strong but not universal. It holds where work *accumulates* against a reality and someone must *answer* for it over time. It weakens or inverts elsewhere, and pretending otherwise would be the kind of map this practice refuses.

- **Throughput work with no accumulating reality.** Tasks valued purely by volume and speed — bulk translation, routine triage, content at scale — have no account to author and no framing for reality to overturn. Here substitution *is* the goal, and this apparatus is dead weight.
- **Disposable generation.** A throwaway prototype, a first-draft brainstorm, a scaffold meant to be deleted: nothing will ever be signed, so there is no authorship to protect. Let the machine author it outright; insisting on the fence would be ceremony.
- **Speed-over-trust domains.** Where fast and approximately right beats slow and accountable, the cost structure that makes the practice worthwhile inverts.
- **Well-modeled, checklist-governed regimes.** Where the domain model has earned authority (axis two, above), follow the map and reserve judgment for deviations. Plan-primacy is not the enemy; *unearned* plan-primacy is.
- **Multi-signer work.** The practice as stated is a solo discipline: one loop, one setpoint, one signer. Nothing here is proven for two comparators sharing traces — whose signature binds, how registered priors compose, what a fence held by a committee is worth. The institutional record is not encouraging: where inspection of machine output has been mandated at scale, it has mostly produced the ritual rather than the inspection (Green, 2022). Scaling the fence is an open problem, stated here as a limit rather than smuggled in as a feature; the transfer test comes first.

The boundary in one line: *the practice applies exactly where the value of the output depends on someone being answerable for it over time, in territory not yet charted well enough for the chart*

to rule.

11 Falsifiability

A manifesto that demands fidelity to reality owes reality the terms of its own surrender. The practice is refuted if any of the following holds at scale and over time:

- **Delegated judgment does not degrade the work.** A practice in which the machine authors the reflection, the scientific decision, or the sign-off — and the resulting body of work proves as trustworthy, deep, and durable as the witnessed version — falsifies the non-delegability claim. The honest test is over time: substitution often looks free until the account must be returned to, defended, or built upon.
- **Opacity beats legibility.** If opaque, database-backed systems reliably outperform plain-text trace architectures on both longevity and human-machine reasoning quality, the substrate commitment is wrong.
- **Hand-groomed maps win at scale.** If maintained prescriptive maps systematically outperform regenerated views once regeneration is cheap, the disposability commitment is mistaken rather than principled.
- **Inspection is ritual.** If practitioners who sign whatever the model emits show no loss in outcome quality, then "responsibility draws the line" describes a ceremony, not a mechanism.

To keep these testable rather than rhetorical, the practice proposes its own measures:

- **Resumption time**: how long to regain full working context on a project untouched for three months — the loop claims this shrinks dramatically.
- **Phantom count**: requirements discovered to trace to no world-trace, per project, with and without the provenance discipline.
- **Signed versus rubber-stamped**: defect and reversal rates of decisions that went through the registered-prior-plus-inspection path versus decisions that did not.
- **Seeded-error detection**: the rate at which deliberately defective machine output, placed blind among the genuine, is caught at inspection — the one unambiguous measure of whether signing is inspection or ceremony.
- **Correction rate over time**: how often the comparator corrects the machine's output, tracked in the knowledge that a decline is ambiguous between model improvement and comparator drift — to be disambiguated (seeded errors are the disambiguator), never read as a success by itself; and that a high rate on a task recorded as delegated is ambiguous between machine failure and Clever Hans — the judge doing the work through the machine — to be disambiguated by the sound check (does the output pass without the judge correcting it, and without the asking having carried the answer?), never read as delegation.
- **The transfer test**: whether the practice survives adoption by a second practitioner without its originator's coaching.

And one limitation stated plainly, because the practice demands it: the evidence base to date is a handful of projects by a single practitioner. Convergence across one person's projects is excellent *hypothesis generation* and weak *confirmation* — a stable personal style would produce the same convergence. The manifesto is offered as a falsifiable proposal, not a proven result. The transfer test is the next experiment.

Since the first draft, external evidence has begun to arrive, and honesty requires logging it on both sides. On the cost of delegated cognition: a survey of 319 knowledge workers finds that the more a worker trusts the model, the less critical-thinking effort they report investing (Lee et al., CHI 2025); an EEG study of LLM-assisted essay writing finds accumulating "cognitive debt,"

with assisted writers unable to quote from their own essays and reporting impaired ownership of them (Kosmyna et al., 2025) — the marathon medal, measured. On the fence as ritual: the empirical record on mandated human oversight of algorithms finds that people largely cannot perform it, and that policy requiring it mostly produces the ceremony (Green, 2022). None of this settles the falsifiers; all of it bears on them, and in the direction the practice predicts — while also documenting how often real practitioners fail to hold the fence.

12 Lineage

The practice is a synthesis, not an invention; naming the ancestry situates it and clarifies the one new thing. The inclusion test for this list is strict: a source earns its place by *doing work* — generating a distinction, constraining a decision, or naming a failure mode — not by resonating. (That test is the practice applied to its own bibliography: an undiscriminated everything-connects reading list is itself an ungoverned reinforcing loop.)

12.1 The engine: representations answerable to reality

- **Alfred Korzybski** (*Science and Sanity*, 1933) — "the map is not the territory"; abstraction loses information at every level. The literal statement of Principle I, applied here recursively: the trace is not the territory either, only the least-derived map of it.
- **Karl Popper** (*Conjectures and Refutations*, 1963) — the loop as epistemology: knowledge grows by exposing representations to refuting contact with reality, never by protecting them. The session-opening re-derivation is conjecture-and-refutation as a habit; the phantom deliverable is a Lakatosian degenerating programme caught in the act. (David Deutsch's "good explanations are hard to vary" is the modern restatement.)
- **W. Ross Ashby** (*Design for a Brain*, 1952) — the law of requisite variety, used here *correctly*: it supplies the criterion the comparator must meet, not a theorem that the comparator must be human.
- **Gregory Bateson** (*Steps to an Ecology of Mind*, 1972) — "information is a difference that makes a difference" is what the comparator detects; circular causation and the map/territory relation are the loop's backbone.
- **Donella Meadows** (*Thinking in Systems*, 2008) — the working vocabulary: stocks, flows, setpoints, delays, leverage points; the insight that the practice's rules sit at the highest-leverage points (the goal, the paradigm, the information flows).
- **John Boyd** — the OODA loop (observe–orient–decide–act): the daily rhythm of this practice is a close cousin, independently derived in the military register, with "orient" as the pivotal and most misunderstood step. A practice-tradition confirmation that re-contacting reality faster than your representation decays is the whole game.
- **W. Edwards Deming and the Toyota tradition** — plan-do-check-act, *genchi genbutsu* ("go and see for yourself at the actual place"), the gemba walk. The industrial orthopraxy of re-contacting reality, with decades of at-scale evidence — the strongest external validation that the loop works outside one practitioner's projects.
- **Event sourcing** (software architecture) — the engineering tradition that *is* this substrate: an append-only event log as the source of truth, all state as disposable projections rebuilt from the log. It brings mature, hard-won failure-mode knowledge (snapshot drift, replay cost) and makes the whole architecture legible to an engineer in one sentence.
- **Architecture Decision Records** (Nygard, 2011) — the engineering orthopraxy of the decision ledger: lightweight records binding each decision to its context, alternatives, and consequences, superseded rather than edited. A closer ancestor of the roadmap's ledger stratum than event sourcing itself.
- **Jay Kreps** ("The Log," 2013) and **Yohei Nakajima** ("The Log is the Agent," 2026) —

the log generalised, then applied to the machines themselves: an ordered, append-only log as the backbone from which every other representation is derived; and the agent-engineering world independently converging on that substrate for language-model agents — log as source of truth, state as disposable projection, end-to-end lineage — with no equivalent of Principle II ~~anywhere in sight~~. *(Narrowed 2026-10-03: human nodes are now in sight in that world, but none is this one — the verifier of Gilda & Gilda (2026) is external to the generation loop and explicitly need not stay human; the industrial human-in-the-loop of 2026 allocates the human by the stakes of the action; and the non-delegability argued in the philosophy of technology (Torrecilla-Pinero, 2026) is ontological and scoped to dignity-touching domains. None allocates by authorship with a fence that moves.)* The clearest external proof that the two principles are separable, and that the second is this practice's actual claim.

- **James C. Scott** (*Seeing Like a State*, 1998) — the cautionary corpus: imposed top-down legibility schemes that flatten reality and break against it. What happens at civilizational scale when the map is allowed to outrank the territory.
- **Carlo Ginzburg** ("Clues," 1979) — the evidential paradigm: reconstructing reality from traces, clues, and symptoms, as the hunter, the physician, and the historian do. The trans-disciplinary pedigree of "traces are the source of truth."
- **George Box** ("Science and Statistics," 1976) — "all models are wrong, some are useful": the pragmatist license for the disposable map, in one page.
- **C. S. Peirce** — abduction: inference to the best explanation from the traces. The logical form of what the comparator does; structural, not decorative.

12.2 The judgment node: why one act stays human

- **Aristotle** (*Nicomachean Ethics*) — *praxis* versus *poiesis*: action whose end is in the doing versus production whose end is the product. The allocation policy, twenty-three centuries early: delegate the poiesis, keep the praxis.
- **Michael Polanyi** (*Personal Knowledge*, 1958; *The Tacit Dimension*,
  1. — "we know more than we can tell": why the judgment draws on

  situated knowledge that cannot be fully externalised into any trace.
- **Hubert Dreyfus** (*Mind Over Machine*) — expert judgment as intuitive, embodied, and non-representational; the novice-to-expert ladder as an account of what the comparator's skill *is* and how it forms.
- **Donald Schön** (*The Reflective Practitioner*, 1983) — reflection-in-action: the working theory of the practice's reflective beat and the inspect-before-signing discipline.
- **John Dewey** (*Logic: The Theory of Inquiry*) — knowing-as-doing; inquiry as the active reconstruction of a problematic situation; the demolition of the spectator theory of knowledge. The practice's native philosophical tradition.
- **James J. Gibson** (*The Ecological Approach to Visual Perception*) — direct perception and affordances: the empirical case that agents pick up reality rather than consulting inner maps of it.
- **Christopher Alexander** (*The Timeless Way of Building*) — organic, piecemeal, reality-corrected growth against the imposed master plan: the live emergent document, in the register of making.
- **Nassim Taleb** (*Skin in the Game*; *Antifragile*) — the binding of decision to consequence as the formal core of signing; and a third systems behaviour beyond balancing and reinforcing — gaining from stressors — as the adversarial temperament the practice needs.
- **The Stoic discipline of assent** (Marcus Aurelius, Epictetus) — assent only to the evident: the oldest statement of the sign-off as a practiced discipline rather than a held belief. Kept from a longer contemplative shortlist because it alone names a move the practice actually makes; the rest of that tradition enriches the felt sense of the work but

grounds no distinction, and is honestly classed as resonance, not lineage.

- **Jesús Torrecilla-Pinero** ("Judgment Cannot Be Delegated: A Subject-Preserving Framework for AI Governance," 2026) — the non-delegability claim stated as an ontological limit: moral judgment is constitutively personal, so delegating it does not transfer it but dissolves the subject who judges; and the observation that a human who rubber-stamps model output satisfies the tracing condition of Meaningful Human Control while exercising nothing that deserves the name of judgment — fragility one, from the governance side. Kept for the contrast it sharpens: he draws the line by ontology, for dignity-touching domains; this practice draws it by authorship, for answerable work, and holds it as a policy whose location moves. Same fence, two different reasons for it.

12.3 The augmentation tradition: extending the witness

- **Vannevar Bush** ("As We May Think," 1945) — the Memex: external associative memory as an extension of recall.
- **J. C. R. Licklider** ("Man-Computer Symbiosis," 1960) and **Douglas Engelbart** ("Augmenting Human Intellect," 1962) — the canonical statements of extend-don't-substitute; this manifesto is their thesis re-derived for the language-model era.
- **Edwin Hutchins** (*Cognition in the Wild*, 1995) and **Clark & Chalmers** ("The Extended Mind," 1998) — cognition lives in artifacts: the checkpoint document is not a record of thinking but a part of it.
- **Endel Tulving** (encoding specificity, 1973) and **Godden & Baddeley** (context-dependent recall, 1975) — the memory science that the retrieval design is *modeled on*: recall works best when retrieval cues match encoding cues. An inspiration honestly labeled as analogy, not implementation.

12.4 The loyal opposition: kept because they push back

- **Judea Pearl** (*The Book of Why*, 2018) — the sharpest scope-limit: the loop self-corrects against *observable* reality, but causal structure is not always observable. Where drift is causal-structural, re-contacting traces will not reveal it.
- **Hayden White** (*Metahistory*, 1973) — the prosecutor of the fidelity claim: every narrative emplots its facts; the "faithful, neutral account" is always already a structured argument. Read against the chronicle, not for it.
- **Kahneman & Klein** ("Conditions for Intuitive Expertise," 2009) — the constraint on the human node: tacit judgment is trustworthy only in environments with valid cues and fast feedback. Both a warning and, read closely, the strongest argument *for* the loop: it is what creates the conditions under which the comparator's intuition can be trusted at all.
- **Atul Gawande** (*The Checklist Manifesto*, 2009) — the evidence that in well-modeled regimes, standing prescriptive maps beat situated judgment. The counter-tradition that forces axis two into the practice instead of leaving it as an unexamined blind spot.
- **Lisanne Bainbridge** ("Ironies of Automation," 1983) — the classic statement that automation makes the human's remaining role harder, not easier: the operator removed from routine work loses exactly the situational awareness the exceptional cases demand, and the better the automation, the less informative the experience of supervising it. The structural cause behind fragility five, four decades before language models.
- **Ethan Mollick** (*Co-Intelligence*, 2024) — the rival allocation policy: delegate by measured capability, mapping the "jagged frontier" empirically and handing off whatever falls inside it. Kept because it names the live alternative this practice refuses: capability-allocation moves the fence wherever the machine improves, with no principled stopping point; authorship-allocation predicts exactly that drift and names what must not cross regardless of capability.

- **Subbarao Kambhampati** ("Can Large Language Models Reason and Plan?", 2024; "LLMs Can't Plan, But Can Help Planning in LLM-Modulo Frameworks", with others, 2024) — the generate-test-critique loop with the human deliberately kept out of the inner iteration: placed at specification and model acquisition, with a sound external verifier certifying the result. Supplies the Clever Hans reading of the correction rate and the test for whether an act has crossed: who certified the answer, and what the prompts carried. Opposition because it draws the human out of the inner loop where fragility five draws them back in to stay calibrated; the two are reconciled by nesting, his loop inside the delegable layer, this one around it.
- **Sankalp Gilda and Shlok Gilda** ("AI-Assisted Engineering Should Track the Epistemic Status and Temporal Validity of Architectural Decisions," 2026) — the oracle rule formalised from the engineering side: every machine-proposed decision carries an epistemic level (conjecture, substantiated, corroborated) and an evidence expiry, and their "Transformer Mandate" requires that whoever finalises a decision be external to the generation loop. Opposition because the mandate is explicitly about separation, not about the human: "currently, this necessitates a human; in future multi-agent systems, the role could theoretically be filled by an independent verifier agent." What moves in both is the fence; what differs is the anchor. Their human stands at ratification until a machine can verify soundly, and then steps aside: allocation by capability, in Kambhampati's shape. This practice moves the line between poiesis and praxis as machines improve, but pins the choice and the signature to answerability and to the value of the doing, which no capability unlocks. Their evidence decay is a clock standing in for re-contact: a decision expires after a set time and is re-validated. Here a decision never expires by the calendar; it is overturned the moment a changed world-trace reaches it, because it is a derived representation like any map and the traces win. The clock is the right device where re-contact is too costly to run on every pass; where it is cheap, re-derivation subsumes it.
- **Human-in-the-loop by stakes** (the 2026 agent-platform pattern: "agent proposes, a person disposes," approval gated on the risk or irreversibility of the action, every intervention logged) — the third allocation policy, beside authorship and capability: the human is placed where the consequences are large, not where the act is theirs. Kept because it is the one that ships, and because it is the arrangement Green (2022) found produces the ritual: a checkpoint chosen by stakes invites a signature chosen by habit.

What is genuinely new here is none of the elements — it is their convergence onto one substrate at one moment: legible digital traces that are simultaneously the practitioner's durable thinking, the machine's working context, and the work's authoritative state, made practicable by machines finally able to read, distil, and narrate traces at negligible cost.

13 Genus and name

13.1 A practice, not a methodology

The right genus for this thing is **practice**, and the precise term is **orthopraxy**. The orthodoxy/orthopraxy distinction is about where authority lives: orthodoxy locates correctness in right belief — a doctrine you assent to; orthopraxy locates it in right practice — conduct done rightly, with doctrine downstream or absent.

Run the practice through that lens. There is no authoritative doctrine to conform to — the map is explicitly disposable, refused as a source of truth; that is the anti-creed move. Authority lives in the doing and its residue: in traces (what was actually done) and in the repeated discipline (re-contact, re-derive, judge, sign). What makes an account genuine rather than counterfeit is not assent to a proposition but whether the witnessing act was actually performed. And the isomorphism is exact:

orthodoxy : orthopraxy :: plan-first : reality-first

Orthodoxy lets doctrine generate conduct; plan-first work lets the plan generate action. The practice refuses both inversions in the same gesture.

One refinement keeps the term honest: this is not a behaviorist orthopraxy of outward conduct. The act it protects is an *interior, situated judgment* — so the precise name for the genus is an **orthopraxy of judgment**: authority lives in practice rather than in any doctrine, but the practice that carries the authority is an act of discernment with an irreducible inside.

13.2 The name

The reality-first practice. The name foregrounds the timeless half — the direction of authority between reality and its representations — and lets the era-dependent half (where exactly the automation fence sits this decade) remain what it is: a policy, revised as machines change, inside a frame that does not.

Digital traces are the practice's substrate; automation is its accelerant; neither is its principle. The principle is the explorer's: the chart is always redrawn from the coastline.

14 The extraction

One passage to carry into any new project:

> Build the work around plain-text, append-only traces that are simultaneously your durable thinking, the machine's working context, and the project's authoritative state. Treat every representation built on them as disposable: re-derived from the traces on every contact, and overruled by the traces wherever they disagree. Hand the machine every act whose value is independent of who performed it — surfacing, orienting, distilling, narrating, scaffolding, and the bulk of the visible labour — because that friction is what otherwise kills the practice. Keep exactly one node: the situated judgment that turns surfaced reality into a committed choice, registered before the machine speaks and signed after adversarial inspection. Hold the fence as a norm and a policy, not a theorem: its location moves as machines improve; the direction of authority — reality over representation — does not. Set the artifact's prescriptiveness by the shape of the work: a live, fused, re-derived roadmap for bounded delivery; a pure append-only spine for open-ended accumulation; an earned checklist where the domain is charted to death. And measure the practice by what it claims: faster resumption, fewer phantoms, an account that can be defended years later by the one who signed it.

15 Provenance of this manifesto

This manifesto distils a longer comparative essay that analyzed the convergent design of the originating projects, together with a critical review of that essay. The review's corrections are applied here rather than appended: the two candidate first principles are presented as a conjunction instead of ranked (the essay's two halves had crowned each in turn); the requisite-variety argument is stated as a criterion rather than a theorem; the memory-science framing is labeled as the analogy it is; the convergence evidence is claimed at honest strength and its single-practitioner limit stated; the anchoring fragility and the plan-primacy counter-tradition are added; and the lineage was re-tiered under an explicit inclusion test, with the additions that test demanded (Aristotle, Popper, Boyd, Deming, event sourcing, Kahneman & Klein, Gawande) and the prunings it forced.

The document was first signed on 2026-06-07. A revision of 2026-08-30 integrated two further sources. A working conversation on oracle use and on the calibration of the judge supplied the oracle rule, the fifth fragility, and the second balancing loop pointed at the comparator. A survey of adjacent work supplied the Bainbridge entry, the rival capability-allocation policy (Mollick), the ADR and log-as-agent lineage, the first external evidence now logged under Falsifiability (Lee et al. 2025, Kosmyna et al. 2025, Green 2022), and the multi-signer scope limit. One sentence — the babysitting-tax claim — was corrected by dated strikethrough rather than rewritten, per the practice's own rule.

A revision of 2026-09-27 added the Clever Hans reading to fragility five and to the correction-rate measure, and a lineage entry, from Subbarao Kambhampati's 2024 essay "Can Large Language Models Reason and Plan?" (Annals of the New York Academy of Sciences), which carries the argument the LLM-Modulo paper of the same year states in one clause: the manifesto had read the correction rate only at its low end, and so named the machine's work passing as the judge's but never the judge's work passing as the machine's. A second reading of the source, before signing, moved the crossing test from the timing of the judge's contribution to its content.

A revision of 2026-10-03 narrowed one lineage claim by dated strikethrough — that the log-as-substrate world had no equivalent of Principle II "anywhere in sight" — after a search found human nodes in that world, none allocated by authorship, and added the three sources that showed it: Torrecilla-Pinero (2026) under the judgment node, Gilda & Gilda (2026) and the human-in-the-loop-by-stakes pattern under the loyal opposition. The claim itself stood; its reach did not.

Per the practice itself, this document is a derived view, not a source of truth. It is downstream of, and disposable relative to, the lived discipline it describes. If the manifesto and the practice ever disagree, the practice wins and the manifesto is re-derived — never the reverse. A practice that could not survive its own manifesto being deleted was never a practice; it was a doctrine wearing a document.

This manifesto, the praxis skill that implements it and the praxis CLI are maintained together at https://github.com/dda410/praxis — the place for corrections, for discussion of where the practice applies and where it does not, and for the dated strikethroughs any revision leaves.

This text is licensed under the Creative Commons Attribution 4.0 International licence (CC BY 4.0); the full licence is in the LICENSE file beside it. Share it, adapt it, build on it, with credit, marking what you changed, and without implying the author endorses the result.